

Exercise 1

Step 1: See a simple LogStash-Pipeline in Action

1. `docker compose -f docker-compose.yml up -d`

2. `docker exec -it elk-stack-dock-logstash-1 /bin/bash`

3. If you start a shell with Docker Desktop enter `bash` to use a more comfortable shell.

4. Find out where logstash is installed

```
which logstash
```

```
/usr/share/logstash/bin/logstash
```

5. `cd` to bin directory of logstash

```
cd /usr/share/logstash/
```

6. Test a simple pipeline with input from and output to commandline.

```
bin/logstash -e 'input { stdin { } } output { stdout { } }'
```

7. When using a docker container with an running Instance of logstash you have to use another `data` directory.

8. You have to start the second logstash process with the flag `--path.data data2`

```
bin/logstash -e 'input { stdin { } } output { stdout { } }' --path.data data2
```

9. Wait until pipeline is running

```
[INFO ] 2022-02-14 07:56:05.167 [Agent thread] agent - Pipelines running
{:count=>1, :running_pipelines=>[:main], :non_running_pipelines=>[]} Then type hello
world
```

```
[2023-03-12T13:35:40,550][INFO ][logstash.javapipeline] [[main] Pipeline started {"pipeline.id"=>"main"}
The stdin plugin is now waiting for input:
[2023-03-12T13:35:40,583][INFO ][logstash.agent] [Pipelines running {:count=>1, :running_pipelines=>[:main], :non_running_pipelines=>[]}
hello world
{
  "@version" => "1",
  "@timestamp" => 2023-03-12T13:35:58.016064900Z,
  "event" => {
    "original" => "hello world"
  },
  "host" => {
    "hostname" => "ad95925a89ce"
  },
  "message" => "hello world"
}
```

10. Stop the pipeline with `STRG-D`

11. List logstash plugins with its version

```
bin/logstash-plugin list --verbose
```

```
logstash@ad95925a89ce:~$ bin/logstash-plugin list --verbose
Using bundled JDK: /usr/share/logstash/jdk
logstash-codec-avro (3.4.0)
logstash-codec-cef (6.2.6)
logstash-codec-collectd (3.1.0)
logstash-codec-dots (3.0.6)
logstash-codec-edn (3.1.0)
logstash-codec-edn_lines (3.1.0)
logstash-codec-es_bulk (3.1.0)
logstash-codec-fluent (3.4.1)
logstash-codec-graphite (3.0.6)
logstash-codec-json (3.1.1)
logstash-codec-json_lines (3.1.0)
logstash-codec-line (3.1.1)
logstash-codec-msgpack (3.1.0)
logstash-codec-multiline (3.1.1)
logstash-codec-netflow (4.3.0)
logstash-codec-plain (3.1.0)
logstash-codec-rubydebug (3.1.0)
logstash-filter-aggregate (2.10.0)
```

12. See if jdbc-integration-plugin is installed

```
bin/logstash-plugin list --verbose | grep 'jdbc'
```

```
logstash@ad95925a89ce:~$ bin/logstash-plugin list --verbose | grep 'jdbc'
logstash-integration-jdbc (5.4.1)
├─ logstash-input-jdbc
├─ logstash-filter-jdbc_streaming
└─ logstash-filter-jdbc_static
```

13. Test the logstash configuration with

```
/usr/share/logstash/bin/logstash --path.settings /usr/share/logstash/config
-t
```

```
Configuration OK
[2023-03-12T13:05:26,789][INFO ][logstash.runner] Using config.test_and_exit mode. Config Validation Result: OK. Exiting Logstash
```

Step 2: Working with LogStach jdbc and PostgreSQL

Work with logstash as shown in [Step 1](#)

The next steps are useful for testing a jdbc-connection to a Postgres container

- replace ip with ip or name of your host
- host.docker.internal:postgresport if you use a docker container with a separate network
- use postgres connector instead of mariadb in the following example

```
bin/logstash --path.data data2 -e 'input {
  jdbc {
    jdbc_connection_string => "jdbc:postgresql://10.3.105.88:5432/onlineshop"
    jdbc_user => "postgres"
    jdbc_password => "geheim"
    jdbc_driver_class => "org.postgresql.Driver"
    statement => "SELECT * FROM visit"
  }
} output { stdout {} }'
```

- Open a commandline
- See this [medium blog](#) for how to setup a pipeline with jdbc.

```
curl https://jdbc.postgresql.org/download/postgresql-42.5.4.jar -o
/usr/share/logstash/logstash-core/lib/jars/postgresql-jdbc.jar
```

1. Create a table in postgres. For Example: orders, visits, order_items

```
psql -U postgres

CREATE DATABASE onlineshop;
\connect onlineshop
```

2. See [Examples](#).
3. create a configuration file to get data from postgres public.orders|visits|order_items in
/usr/share/logstash/pipeline/logstash.conf
4. See [jdbc-input Doku](#)
5. Add tracking for PK column with :sql_last_value
6. You have to add a line for scheduling to make logstash run without interruption: `schedule => "*/5 * * * *`
 - This file is mapped to `<path-to-your-docker>/elk-stack-dock/logstash/pipeline/logstash.conf`
 - optional: configure a pipeline in `/usr/share/logstash/config/pipelines.yml` (a default is there). This file is not mapped.

Complete file:

```
input {
  jdbc {
    jdbc_connection_string => "jdbc:postgresql://10.3.105.88:5432/onlineshop"
    jdbc_user => "postgres"
    jdbc_password => "geheim"
    jdbc_driver_class => "org.postgresql.Driver"
    statement => "SELECT * FROM orders WHERE idorders > :sql_last_value"
    use_column_value => true
    tracking_column => "idorders"
    schedule => "*/5 * * * *"
  }
}
```

```

    }
  }

  output {
    elasticsearch {
      hosts => "elasticsearch:9200"
      user => "logstash_internal"
      password => "${LOGSTASH_INTERNAL_PASSWORD}"
      index => "orders"
      document_id => "orders_%{idorders}"
      doc_as_upsert => true
    }
  }
}

```

Execute pipeline with `bin/logstash -f pipeline/logstash.conf --path.data data2`

```

[2023-03-12T13:28:32,663][INFO ][logstash.javapipeline] [main] Pipeline started {"pipeline.id"=>"main"}
[2023-03-12T13:28:32,698][INFO ][logstash.agent] Pipelines running {:count=>1, :running_pipelines=>[:main], :non_running_pipelines=>[]}
[2023-03-12T13:30:00,891][INFO ][logstash.inputs.jdbc] [main][d79e5918a1cfd9845ffc05bd30e2e91e18daac4b244978af98fd4d4b48cbc8ba] (0.013221s) SELECT * FROM orders WHERE idorders > 45

```

7. Go to Kibana -> Management -> Dev Tools and query logstash index, to see the data rows stored in ES by the pipeline built.

The screenshot shows the Kibana Dev Tools interface. On the left, the 'Console' tab is active, displaying a search query: `GET _search` with a body containing `"query": { "match_all": {} }`. On the right, the search results are displayed in a JSON format. The results show two documents from the `.ds-logs-generic-default-2023.03.01-000001` index. Each document has a unique `_id`, a score of 1, and a source object containing log data such as `ip_address`, `idvisit`, `timestamp`, `user_iduser`, `@timestamp`, `@version`, `data_stream`, `type`, `dataset`, and `namespace`.

8. Compare with MongoDB Pipeline

- See <https://www.mongodb.com/developer/products/mongodb/mongoimport-guide/> to get a glue how MongoDB supports you in this case.
- See for a video-description: <https://www.mongodb.com/developer/videos/import-data-into-mongodb-3-ways---bonus---export-data-from-postgres/>
- MongoDB Compass can be installed locally, but Docker project was last updated in 2021.

9. Write a short summary, what is different to ELK and give a statement, what you think is easier to automate.

Both are tools for managing and analyzing data, but they have different approaches.

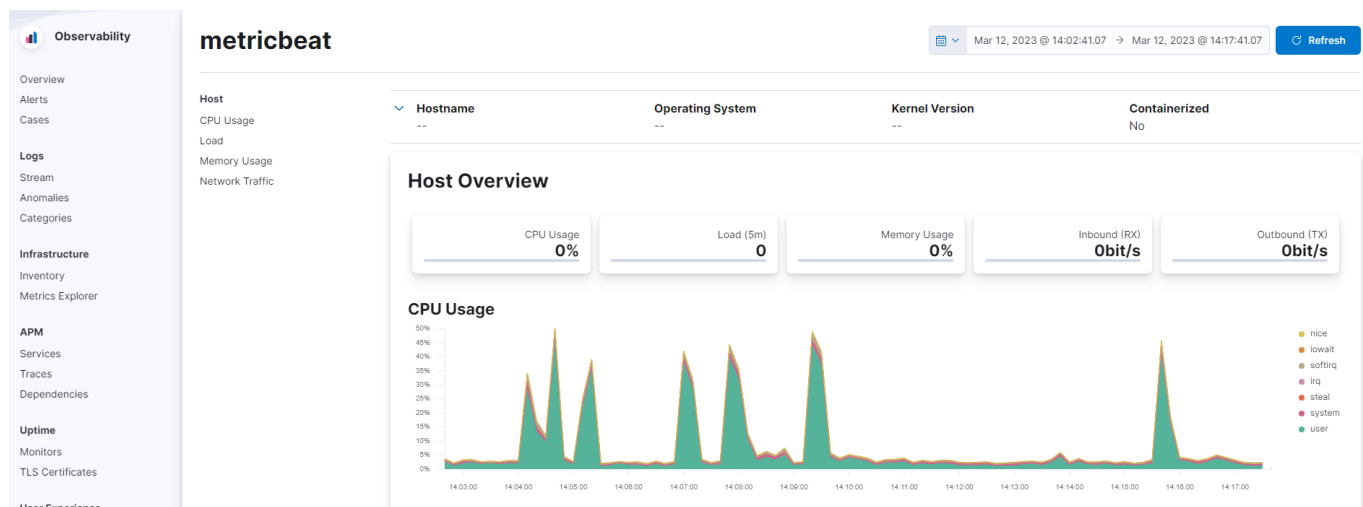
The elk stack is designed for real-time analysis of log data. MongoDB pipeline is a tool for manipulating and processing data in MongoDB, which is a NoSQL document based database. It is designed for more general data processing. For complex data MongoDB would be easier because it provides a wide range of functions and operations with the Aggregation framework or MapReduce.

Step 3: Adding a Metricbeat pipeline and build dashboards in Kibana

- See [Metricbeat README](#) for this.
- First only start the ELK-Stack. See logs of Kibana, to see if all services are already running.
- In a browser connect to <http://localhost:5601>. username/pwd: elastic/changeme
- Then use the command to start metricbeat container

```
docker-compose -f docker-compose.yml -f extensions/metricbeat/metricbeat-compose.yml up
```

1. activate requested passwords for metricbeat in `<path-to-your-docker>/elk-stack-dock/.env`.
2. rerun setup container
3. restart metricbeat



4. Test, if data are retrieved by elasticsearch

- `curl -u elastic:changeme http://localhost:9200/_search`
- `curl -u elastic:changeme http://localhost:9200/logstash/_search`
- `curl -u elastic:changeme http://localhost:9200/logstash/_search?pretty=true`

```

PS C:\Users\Andi\Documents\Github\elk-stack-dock> curl -u elastic:changeme http://localhost:9200/orders/_search?pretty=true
{
  "took" : 0,
  "timed_out" : false,
  "_shards" : {
    "total" : 1,
    "successful" : 1,
    "skipped" : 0,
    "failed" : 0
  },
  "hits" : {
    "total" : {
      "value" : 5,
      "relation" : "eq"
    },
    "max_score" : 1.0,
    "hits" : [
      {
        "_index" : "orders",
        "_id" : " orders_15 ",
        "_score" : 1.0,
        "_source" : {
          "date_ordered" : "22.02.2024",
          "@version" : "1",
          "idorders" : 15,
          "@timestamp" : "2023-03-11T10:50:00.587772500Z",
          "married_to" : null,
          "user_iduser" : 2,
          "total_sum" : 4000.0
        }
      },
      {
        "_index" : "orders",
        "_id" : " orders_20 ",
        "_score" : 1.0,
        "_source" : {
          "date_ordered" : "22.02.2024",
          "@version" : "1",
          "idorders" : 20,
          "@timestamp" : "2023-03-11T10:50:00.588913900Z",
          "married_to" : null,
          "user_iduser" : 2,
          "total_sum" : 8000.0
        }
      },
      {
        "_index" : "orders",
        "_id" : " orders_21 ",
        "_score" : 1.0,
        "_source" : {
          "date_ordered" : "22.02.2024",
          "@version" : "1",
          "idorders" : 21,
          "@timestamp" : "2023-03-11T10:55:00.032829600Z",
          "married_to" : null,
          "user_iduser" : 2,
          "total_sum" : 100.0
        }
      }
    ]
  }
}

```

See [Query DSL Doku]

(<https://www.elastic.co/guide/en/elasticsearch/reference/8.6/query-dsl.html>) for the next steps.

5. either find the curl command to query for a specific entry in a table with a match

```
curl -u elastic:changeme "http://localhost:9200/orders/_doc/orders_30?pretty"
```

```
PS C:\Users\Andi\Documents\Github\elk-stack-dock> curl -u elastic:changeme "http://localhost:9200/orders/_doc/orders_30?pretty"
{
  "_index" : "orders",
  "_id" : "orders_30",
  "_version" : 1,
  "_seq_no" : 5,
  "_primary_term" : 1,
  "found" : true,
  "_source" : {
    "user_iduser" : 2,
    "idorders" : 30,
    "date_ordered" : "22.02.2024",
    "total_sum" : 1000.0,
    "@version" : "1",
    "@timestamp" : "2023-03-11T11:20:00.733387100Z",
    "married_to" : null
  }
}
```

6. and one, that doesn't match any database entry

```
PS C:\Users\Andi\Documents\Github\elk-stack-dock> curl -u elastic:changeme "http://localhost:9200/orders/orders_35?pretty"
{
  "error" : "Incorrect HTTP method for uri [/orders/orders_35?pretty] and method [GET], allowed: [POST]",
  "status" : 405
}
```

7. Connect to Kibana in Browser

8. Goto Burger Menu -> Management -> Dev Tools

9. Query for all entries in logstash index.

The screenshot shows the Kibana Dev Tools interface. On the left, the 'GET orders' query is entered. On the right, the mapping for the 'orders' index is displayed, showing fields like '@timestamp', '@version', 'date_ordered', 'idorders', 'total_sum', and 'user_iduser' with their respective types and mappings.

11. Query for a specific entry, like you did with curl

GET orders/_doc/orders_25

The screenshot shows a web application interface with a top navigation bar containing 'Console', 'Search Profiler', 'Grok Debugger', and 'Painless Lab' (with a 'BETA' badge). Below the navigation bar, there are tabs for 'History', 'Settings', 'Variables', and 'Help'. On the right side of the interface, there are two status boxes: a green one saying '200 - OK' and a grey one showing '25 ms'. The main area is split into two panes. The left pane contains a REST client interface with a text area showing a cURL command: `1 # Click the Variables button, above, to create your own variables.`, `2 GET orders/_doc/%20orders_25%20`, and `3`. The right pane displays a JSON response from a REST client, with line numbers 1 through 17 on the left. The JSON object is:

```
{
  "_index": "orders",
  "_id": " orders_25 ",
  "_version": 1,
  "_seq_no": 2,
  "_primary_term": 1,
  "found": true,
  "_source": {
    "date_ordered": "22.02.2024",
    "@version": "1",
    "idorders": 25,
    "@timestamp": "2023-03-11T10:55:00.033413Z",
    "married_to": null,
    "user_iduser": 2,
    "total_sum": 1000
  }
}
```

12. Sort the entries in asc and desc order

Query all and sort with total_sum

```
GET orders/_search
{
  "sort" : [
    { "total_sum": { "order": "desc" } }
  ],
  "query": {
    "match_all": {}
  }
}
```